

Explainable Deep Learning Framework for Automated Diabetic Retinopathy Screening

A Report on Major Project Part 1
submitted in partial fulfillment of the requirements for
the award of the degree of

Bachelor of Engineering

in

Artificial Intelligence and Data Science

By

Malige Navya (1601-22-771-080)

Mohammed Danish Ahmed (1601-22-771-116)

Umar Uz Zama Mohammed (1601-22-771-131)

Under the esteemed guidance of

Mr. Naveen Raja SM

Assistant professor



**DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND DATA SCIENCE
CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY
HYDERABAD – 500075**

May 2026



**DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND DATA SCIENCE
CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY
HYDERABAD – 500075**

INSTITUTE VISION

“To be the center of excellence in technical education and research”.

INSTITUTE MISSION


“To address the emerging needs through quality technical education and advanced research”.

DEPARTMENT VISION

”To be a globally recognized center of excellence in the field of Artificial Intelligence and Data Science that produces innovative pioneers and research experts capable of addressing complex real-world challenges and contributing to the socio-economic development of the nation.”

DEPARTMENT MISSION

1. To provide cutting-edge education in the field of Artificial Intelligence and Data Science that is rooted in ethical and moral values.
2. To establish strong partnerships with industries and research organizations in the field of Artificial Intelligence and Data Science, and to excel in the emerging areas of research by creating innovative solutions.
3. To cultivate a strong sense of social responsibility among students, fostering their inclination to utilize their knowledge and skills for the betterment of society.
4. To motivate and mentor students to become trailblazers in Artificial Intelligence and Data Science, and develop an entrepreneurial mindset that nurtures innovation and creativity.



**DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND DATA SCIENCE
CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY
HYDERABAD – 500075**

DECLARATION CERTIFICATE

We hereby declare that the project titled **Explainable Deep Learning Framework for Automated Diabetic Retinopathy Screening** submitted by us to the **Artificial Intelligence and Data Science CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY, HYDERABAD** in partial fulfillment of the requirements for the award of **Bachelor of Engineering** is a bona-fide record of the work carried out by us under the supervision of **Mr. Naveen Raja SM** . We further declare that the work reported in this project, has not been submitted and will not be submitted, either in part or in full, for the award of any other degree or diploma of this institute or of any other institute or University.

Project Associates

Malige Navya (1601-22-771-080)

Mohammed Danish Ahmed (1601-22-771-116)

Umar Uz Zama Mohammed (1601-22-771-131)



**DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND DATA SCIENCE
CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY
HYDERABAD – 500075**

ANTI-PLAGIARISM DECLARATION

We hereby declare that the project titled **”Explainable Deep Learning Framework for Automated Diabetic Retinopathy Screening”** is a record of original work carried out by us. The report has been verified for originality using the **Turnitin** software on May 4, 2026, which revealed a similarity index within the acceptable range of **20%**. This result is well within the permissible limits prescribed by the Institute.

We further declare that all sources of information and literature used in this report have been duly acknowledged. We are aware that any act of plagiarism will lead to disciplinary action as per the rules and regulations of the Institute.

Date: May 4, 2026

Place: Hyderabad

Project Associates

Malige Navya (1601-22-771-080)

Mohammed Danish Ahmed (1601-22-771-116)

Umar Uz Zama Mohammed (1601-22-771-131)

Supervisor Signature

Mr. Naveen Raja SM



DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND DATA SCIENCE
CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY
HYDERABAD – 500075
BONAFIDE CERTIFICATE

This is to certify that the project titled **Explainable Deep Learning Framework for Automated Diabetic Retinopathy Screening** is a bonafide record of the work done by

Malige Navya (1601-22-771-080)

Mohammed Danish Ahmed (1601-22-771-116)

Umar Uz Zama Mohammed (1601-22-771-131)

in partial fulfillment of the requirements for the award of the degree of **Bachelor of Engineering in Artificial Intelligence and Data Science** to the **CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY, HYDERABAD** carried out under my guidance and supervision during the year 2024-25. The results presented in this project report have not been submitted to any other university or Institute for the award of any degree.

Mr. Naveen Raja SM

Guide

Dr. K. Radhika

Head of the Department

Submitted for Semester Major-Project Part 1 viva-voce examination held on _____

Examiner-1

Examiner-2

SUSTAINABLE DEVELOPMENT GOALS (SDGs) MAPPED

The proposed Explainable Deep Learning Framework for Automated Diabetic Retinopathy Screening contributes to multiple United Nations Sustainable Development Goals. The system ensures early detection, accessibility, and transparency using explainable AI techniques.



Figure 1: Sustainable Development Goals mapped to the proposed project

SDG	Goal Title	Relevant Target	Project Contribution	System Feature
3	Good Health and Well-Being	Target 3.4: Reduce mortality; Target 3.d: Early detection	Early detection of diabetic retinopathy to prevent blindness	Retinal image classification; automated screening
8	Decent Work and Economic Growth	Target 8.5: Employment; Target 8.6: Reduce workload	Reduces burden on ophthalmologists through automation	Batch screening; automated reports
10	Reduced Inequalities	Target 10.2: Inclusion; Target 10.3: Equal healthcare access	Affordable screening for rural and underserved populations	Web-based system; remote diagnosis
16	Peace, Justice and Strong Institutions	Target 16.6: Transparency; Target 16.10: Access to information	Ensures trust using explainable AI models	Grad-CAM visualization; interpretable outputs

Table 1: Mapping of Project to United Nations Sustainable Development Goals

ABSTRACT

Diabetic retinopathy (DR) is a leading cause of avoidable blindness globally, with 103 million people affected. Early diagnosis minimises the risk of permanent blindness; but many rural and low-resource communities lack access to specialist ophthalmological screening. This project is an enterprise-grade, explainable deep learning solution to automate and provide a clinician-trustworthy screening of diabetic retinopathy in retinal fundus images.

RetinaIQ uses a multi-stage convolutional neural network (CNN) architecture with EfficientNetV2-B3 as the main production model, which is able to grade retinal images into five standard diabetic retinopathy levels (0-4) in accordance to the American Academy of Ophthalmology (AAO) 2023 guidelines and NICE NG28. RetinaIQ employs three separate explainability techniques - Gradient-weighted Class Activation Mapping (Grad-CAM), Local Interpretable Model-Agnostic Explanations (LIME), and SHapley Additive exPlanations (SHAP) - to produce visual and interpretable masks that allow clinicians to verify the AI predictions at the pixel level. The system backend uses FastAPI with asynchronous task queues (Celery and Redis) to provide RESTful API endpoints for authentication, patient, prediction, and PDF report generation. It includes an adaptive image preprocessing pipeline that applies quality-aware Contrast Limited Adaptive Histogram Equalization (CLAHE), bilateral denoising, gamma correction, and optic disc masking to standardize the quality of the input images. A smart model router adaptively routes model inference based on workload, hardware, and image size. The system stores patient information, prediction artifacts and explainability results in PostgreSQL for relational data, and MinIO S3-compatible object storage for image and report management. A web-based single-page application (SPA) built using the React framework offers clinicians an easy way to upload images, view prediction results, view patient timelines, and generate report PDFs. The system includes HIPAA-compliant audit logging, role-based access control, and Json Web Tokens (JWT) for user authentication. In experiments, RetinaIQ attains a training accuracy of 96.1 percentage and an Area Under the Curve (AUC) of 0.987 with EfficientNetV2-B3 model, with an average inference time of 385 milliseconds per image. The entire system is Docker containerized and can run on clinical workstations, cloud systems, and telemedicine systems.

Keywords : Diabetic Retinopathy, Deep Learning, Explainable AI, Grad-CAM, LIME, EfficientNetV2, FastAPI, Medical Image Classification.

ACKNOWLEDGEMENTS

We would like to express our deepest gratitude to the following people for guiding us through this course and without whom this project and the results achieved from it would not have reached completion.

Mr. Naveen Raja SM, Assistant Professor, Department of Artificial Intelligence and Data Science, for helping us and guiding us in the course of this project. Without her guidance, we would not have been able to successfully complete this project. Her patience and genial attitude is and always will be a source of inspiration to us.

Dr. K. Radhika, the Head of the Department, Department of Artificial Intelligence and Data Science, for allowing us to avail the facilities at the department.

We are also thankful to the faculty and staff members of the Department of Artificial Intelligence and Data Science, our individual parents and our friends for their constant support and help.

TABLE OF CONTENTS

Title	Page No.
DECLARATION	i
ANTI-PLAGIARISM DECLARATION	ii
BONAFIDE CERTIFICATE	iii
ABSTRACT	v
ACKNOWLEDGEMENTS	vi
TABLE OF CONTENTS	vii
LIST OF TABLES	x
LIST OF FIGURES	xi
CHAPTER 1 INTRODUCTION	1
1.1 Introduction to the Topic	1
1.1.1 Background: Retinal Fundus Image Analysis	1
1.1.2 Background: Explainable Deep Learning for Medical Imaging . . .	2
1.1.3 Background: Transfer Learning for Healthcare AI	2
1.2 Problem Definition	3
1.2.1 Challenges in Current DR Monitoring	3
1.2.2 Absence of Real-Time Corrective Feedback	3
1.2.3 Limited Dataset Diversity and Annotation Scarcity	4
1.2.4 Problem Statement	4
1.2.5 Objectives and Outcomes	4
1.3 Motivation	6
1.4 Scope of the Project	6
1.5 Thesis Organization	6
CHAPTER 2 LITERATURE REVIEW	8
2.1 Overview	8
2.2 Evolution of Diabetic Retinopathy Detection Systems	8
2.2.1 Phase 1: Heuristic and Rule-Based Systems (2010–2015)	8
2.2.2 Phase 2: Machine Learning Methods (2015–2018)	9
2.2.3 Phase 3: Deep Learning and Explainable AI Methods (2018 – Present)	9
2.3 Related Works	10
2.3.1 Paper 1: A Hybrid Approach for Diagnosing Diabetic Retinopathy from Fundus Image Exploiting Deep Features	10

2.3.2	Paper 2: A Hybrid Approach using customized machine learning classifiers and multiple feature extractors	10
2.3.3	Paper 3: Attention-Based Deep Learning Framework for Automatic Fundus Image Processing	11
2.3.4	Paper 4: Automated Diabetic Retinopathy Detection Using Ensemble Based Classification Model	12
2.4	Summary Table about Literature Survey	13
2.5	Gaps Identified	13
2.5.1	Gap 1: Difficulty to Explain	13
2.5.2	Gap 2: Real-Time Deployment	13
2.5.3	Gap 3: Dataset Diversity and Data Imbalance	13
2.5.4	Gap 4: Transfer Learning and Low-Data Deployment	14
2.5.5	Gap 5: Explainable Clinical Decision Support	14
2.5.6	Gap 6: Privacy-Preserving Healthcare Deployment	14
2.5.7	Gap 7: Healthcare Workflow Integration	14
CHAPTER 3 PROPOSED SYSTEM		15
3.1	Introduction	15
3.2	Methodology	15
3.2.1	Overall Pipeline	15
3.3	Retinal Image Representation	16
3.3.1	Image Formulation	16
3.3.2	Image Normalization	16
3.4	Spatial Feature Extraction	17
3.5	Multi-Scale Deep Learning Framework	17
3.5.1	Construction of Multi-Scale Feature Maps	17
3.5.2	EfficientNetV2-B3 Layer Architecture	18
3.6	Explainability Mechanism	18
3.6.1	Multi-Method Explainable AI Framework	18
3.6.2	Interpretation of Attention from Clinical Perspective	19
3.7	Complete Model Architecture	19
3.8	Dataset Strategy	20
3.8.1	Transfer Learning Deployment	20
3.8.2	Hybrid Multi-Dataset Training	20
3.8.3	Model Performance	21
3.9	System Design Considerations	21
3.9.1	Computational Efficiency	21
3.9.2	Privacy-by-Design	21
3.9.3	Accessibility and Inclusivity	22
3.10	Clinical Validation and DR Severity Classification	22
3.10.1	DR Severity Classification	22
3.10.2	Explainability-Based Clinical Validation	22

3.11	Custom Dataset Collection Protocol	23
3.11.1	Image Collection Setup	23
3.11.2	Subjects and Retinal Conditions	23
3.11.3	Annotation Protocol	23
3.11.4	Workflow Description	24
CHAPTER 4	IMPLEMENTATION	26
4.1	Development Environment	26
4.1.1	Hardware Configuration	26
4.1.2	Software Libraries and Their Versions	26
4.1.3	Project Setup and Environment	27
4.1.4	Data Preprocessing Implementation	27
4.1.5	Feature Extraction and Model Architecture	28
4.1.6	Explainable AI (XAI) Implementation	29
4.1.7	Training and Hyperparameter Configuration	30
4.1.8	Diagnostic Backend API Integration	30
CHAPTER 5	RESULTS AND DISCUSSION	32
5.1	Gantt Chart: Project Timeline	32
5.2	Quantitative Results (Model Performance)	33
5.2.1	Quantitative Metrics	33
5.2.2	Confusion Matrix Analysis	33
5.3	System Deployment & Interface	35
5.3.1	The Home Page	35
5.3.2	Login Page	36
5.3.3	Registration Page	37
5.3.4	Dashboard	38
5.3.5	Image Analysis Page	39
5.3.6	The Patient Management Page	41
5.3.7	Settings Page	42
CHAPTER 6	CONCLUSION AND FUTURE SCOPE	43
6.1	Conclusion	43
6.2	Future Scope	43
CHAPTER 7	BIBLIOGRAPHY	45

LIST OF TABLES

1	Mapping of Project to United Nations Sustainable Development Goals	iv
4.1	Hardware Specifications Used for Development and Testing	26
4.2	Software Requirements and Versions	26
5.1	Quantitative Metrics for DR Diagnostic Model	33
5.2	Confusion Matrix across 5 DR Grades	34

LIST OF FIGURES

1	Sustainable Development Goals mapped to the proposed project	iv
3.1	Workflow for Diabetic Retinopathy Detection	24
4.1	Integrated Development Architecture and Modular Environment Configura- tion for the RetinalIQ Framework	27
4.2	Multi-Stage Image Preprocessing Pipeline for Noise Reduction and Contrast Enhancement of Fundus Images	28
4.3	Hierarchical Feature Extraction Layering utilizing the EfficientNetV2-B3 Backbone for Multi-Class Grading	29
4.4	Gradient-weighted Class Activation Mapping Implementation for Lesion- Level Interpretability	29
4.5	Convergence Optimization through Hyperparameter Tuning and Categorical Cross-Entropy Loss Monitoring	30
4.6	Asynchronous API Architecture for Seamless Telemedicine Integration and Real-Time Diagnostic Inference	31
5.1	Gantt chart	32
5.2	RetinalQ User Landing Interface with Core Capability Highlights and Clin- ical Motivation	36
5.3	Secure Authentication Gateway Featuring Multi-Modal Login and Regula- tory Medical Disclaimer	37
5.4	User Onboarding Interface for HIPAA-Compliant Practitioner Registration .	38
5.5	Authenticated Clinical Dashboard with Real-Time KPIs and Population Severity Distributio	39
5.6	Integrated Diagnostic Interface for Fundus Image Processing and Natural Language Clinical Feedback	40
5.7	Explainability Module Displaying LIME-Based Superpixel Localization for Model Decision Transparency	41
5.8	Patient Management System for Longitudinal Record Tracking and Demo- graphic Data Storage	41
5.9	System Configuration and User Profile Management with Persistent Safety Disclosures	42

LIST OF ABBREVIATIONS

Abbreviation	Full Form
AI	Artificial Intelligence
AUC	Area Under ROC Curve / Area Under the Curve
CBAM	Convolutional Block Attention Module
CBIT	Chaitanya Bharathi Institute of Technology
CLAHE	Contrast Limited Adaptive Histogram Equalization
CNN	Convolutional Neural Network(s)
DL	Deep Learning
DR	Diabetic Retinopathy
GAN	Generative Adversarial Networks
LIME	Local Interpretable Model-Agnostic Explanations
OCT	Optical Coherence Tomography
QWK	Quadratic Weighted Kappa
SE-LALG	Stacked Ensemble-based Logistic Adaboost Light Gradient
SHAP	SHapley Additive exPlanations
SMOTE	Synthetic Minority Over-sampling Technique
ViT	Vision Transformer
XAI	Explainable Artificial Intelligence

CHAPTER 1

INTRODUCTION

1.1 Introduction to the Topic

Diabetic retinopathy (DR) is a serious complication of diabetes mellitus which is one of the most common reasons of preventable blindness throughout the world. It causes clinical changes in the retina, including microaneurysms, retinal hemorrhages, exudates, and macular edema, that occur when high blood glucose levels over a long period damage the small blood vessels in the eye. In many instances, the early stages of the disease do not cause any symptoms, and early detection is critical to avoid permanent vision loss. The traditional diagnostic approach relies on manual examination of retinal fundus images by ophthalmologists, which is time-consuming and labor-intensive. Although accurate, it is difficult to scale to a large population.

As Artificial Intelligence and Deep Learning become increasingly practical, automated analysis of retinal images has gained importance as a tool for mass screening and early detection. Deep learning systems can accurately categorize retinal images into severity levels of DR and support doctors in identifying urgent cases. However, many current models behave as “black boxes,” producing predictions without clearly explaining the reasoning behind them. This limits trust and slows clinical adoption in healthcare environments. Because of this, an explainable deep learning framework is particularly useful, as it can not only identify the presence of DR but also display the visual evidence underlying the prediction. The proposed project aims to develop such a system for reliable, transparent, and deployable DR screening.

1.1.1 Background: Retinal Fundus Image Analysis

Retinal fundus imaging is a non-invasive imaging technique used to examine the interior surface of the eye, including the retina, optic disc, blood vessels, and macula. These images are highly useful for identifying different stages of diabetic retinopathy. However, the challenge lies in the fact that fundus images acquired in real-world settings vary significantly in brightness, contrast, blur, orientation, and overall quality. Poor image quality can conceal lesions and distort vessel patterns, directly affecting classification performance. Due to this, preprocessing becomes an important step in constructing an accurate DR detection system.

In the proposed project, fundus images are standardized using an adaptive preprocessing pipeline before being used for further downstream analysis. Techniques such as contrast enhancement, denoising, and quality-based correction improve the visibility of retinal abnormalities before classification is performed. This is important because the success

of deep learning in medical imaging depends not only on the model architecture but also on the quality and preparation of the input images. A robust preprocessing stage improves system reliability and makes the model more resistant to variations caused by different devices and clinical environments.

1.1.2 Background: Explainable Deep Learning for Medical Imaging

Deep learning has achieved promising results in medical image classification tasks such as disease detection from radiological and retinal images. Convolutional Neural Networks (CNNs) are capable of learning complex visual patterns that may not be easily visible to the human eye, including lesion shapes, vessel abnormalities, and texture differences. Because of this capability, CNN-based models are well suited for automatic diabetic retinopathy screening. In this project, EfficientNetV2-B3 is used as the primary production model along with a multi-stage classification setup.

Medical applications, however, require interpretability. Producing a correct prediction alone is not sufficient unless clinicians are able to understand why the prediction was made. To address this challenge, the project integrates Grad-CAM, LIME, and SHAP to generate explainable outputs that highlight the retinal regions influencing the prediction. These visual explanations assist doctors in verifying whether the model is focusing on clinically meaningful areas, thereby improving confidence in the system and supporting safer deployment in real-world screening scenarios.

1.1.3 Background: Transfer Learning for Healthcare AI

One of the main obstacles to the successful integration of AI in healthcare is the scarcity of extensive, high-quality medical data sets that are properly labeled and structured. Collection and labelling of clinical images often need experts, and this is expensive and time consuming. Transfer learning is one way to tackle this issue by adapting a large model, or a powerful backbone, to a particular medical task. This proves particularly beneficial in diabetic retinopathy (DR) diagnosis, where the amount of available data might not be sufficient for training a deep model from scratch.

The usage of transfer learning in the case of this project is to learn useful image features like edges, textures, lesion-like patterns, and spatial structures from fundus images. This enhances training efficiency, generalization and decreases the amount of data required for task specific training. Transfer learning is particularly useful for healthcare AI, as it enables the creation of more efficient and viable systems that can perform well in real clinical settings with reduced time. In healthcare AI, transfer learning can facilitate the development of quicker, more viable systems that can successfully meet the demands of real clinical settings.

1.2 Problem Definition

The key challenge for this project is to develop a system to detect DR that is accessible, automated, explainable, and deployable in low resource and high-volume healthcare environments. Manual grading of fundus images requires the expertise of an ophthalmologist and is therefore expensive, hard to access and subjective. While automated models are available, many do not offer prediction label and are hard to be trusted by clinicians in practice.

Another key difficulty is that images of the retina from the real world are likely to be noisy and varying. Unless the preprocessing is designed carefully, there may be a decrease in the accuracy of a model due to changes in illumination, focus, contrast, and orientation. Furthermore, a clinical system should be capable of meeting the workflow needs of patients, including recording patient details, predicting patient results, reporting, secure authentication, and audit logging. The challenge, then, is to not only be able to detect diabetic retinopathy, but to also develop a comprehensive intelligent system, accurate, explainable and ready to be used in clinical practice.

1.2.1 Challenges in Current DR Monitoring

There are a number of practical challenges with current DR screening methods. The first problem is that retinal images can only be graded by specialist doctors, making mass screening difficult because there is not enough ophthalmologist availability in these areas. The second challenge is the opacity of the many AI models, as black box predictions lower trust in the AI and aid clinical validation more difficult. The third challenge is that the quality of the retinal images captured in hospitals, mobile screening units, and cameras varies a lot, and impacts the performance of the model, if the systems are not resilient to this variability.

Another issue is deployment complexity. Despite successful experimental results, it can be challenging to incorporate a model into actual clinical practice. A DR screening solution should be scalable, reliable, secure, and support authentication, storage, reporting and processing. These are concerns that make it more than a mere classification problem for images, it's a full end-to-end healthcare automation problem.

1.2.2 Absence of Real-Time Corrective Feedback

The equivalent of real time corrective feedback in diabetic retinopathy screening is prompt decision support. A screening system cannot just generate predictions after a long delay, but it should be used by the clinician or health workers to identify cases for review, quickly. A slow response time diminishes usefulness in busy clinics and screening camps that require a quick response for efficiency.

This is what your project is going to do: Design a production ready backend using Fast API, asynchronous processing and model routing to enable faster inference and reporting.

This makes the system more applicable to real-world clinical and multi-user applications such as telemedicine, screening programs, etc. where turnaround time is important.

1.2.3 Limited Dataset Diversity and Annotation Scarcity

For your project, this section should be modified to account for limited diversity in your DR data sets and limited annotations. Retinal images typically are not representative of the diversity of image quality, disease presentation, camera types and patient demographics encountered in real practice, particularly when acquired from the public domain. This leaves it hard to train a model with a small amount of data to give a good generalisation across hospitals and across screening settings.

Still another issue is the imbalance between classes, particularly if early stage disease classes are under-represented. Mild and moderate cases are clinically relevant and significant for early intervention, and the system needs to be trained carefully so that it does not overfit to the majority class. That's why it's crucial you have a strong pre-processing and model design strategy for your project.

1.2.4 Problem Statement

The system should, given one of a retinal fundus images from a diabetic patient:

1. preprocess the image to make the retina more visible and reduce the noise,
2. extract meaningful visual features using a deep learning model,
3. label the image with the severity of DR,
4. produce interpretable visual results that identify key retinal areas,
5. make the prediction and report within a realistic time period for clinical practice,
6. securely store the outcome for patient tracking and follow-up.

The system needs to accomplish these goals, and to be accurate, explainable, secure, and deployable in real healthcare scenarios.

1.2.5 Objectives and Outcomes

Primary Objectives

The proposed project aims to:

1. To create an automated deep learning system that can classify diabetic retinopathy (DR) from retinal fundus images into clinically meaningful DR severity categories.
2. Integrating explainability methods like Grad-CAM, LIME and SHAP to enable the system to offer visual and interpretable explanations for its predictions.

3. To create an adaptive preprocessing pipeline which will enhance image quality and enable strong model performance under different input conditions.
4. To create a production-ready backend to support secure data for predictions, patient management and report production.
5. To facilitate the use of images in clinical practice via a web-based application for uploading images, viewing results, and tracking history of screening.
6. To deploy in a secure manner with authentication, access control, audit logging and orderly data storage.

Secondary Objectives

The secondary objectives are:

1. To make the predictions of deep learning more transparent for medical imaging.
2. To overcome the need for manual grading of fundus images by an ophthalmologist in routine screening.
3. To adapt the system to be deployable in Telemedicine and Low-resource environments.
4. To help make rapid decisions for effective backend processing and reporting.
5. To provide a structure for expanding to other ophthalmic diseases and/or medical imaging tasks.

Expected Outcomes

The following are the project expected outcomes:

1. A successful diabetic retinopathy screening system that has excellent classification ability.
2. Explanations with significance that demonstrate the region of the retina that affects the model's prediction.
3. Greater confidence in AI diagnostics with visual feedback.
4. Easier and more efficient clinical screening help.
5. A deployable architecture that can be used in clinical and telemedicine environments.

1.3 Motivation

The need for this project is due to the ever growing burden of diabetes in the world and the need to detect DR at an early stage to prevent irreversible vision loss. DR can develop without any early signs or symptoms, so screening is an important component. Manual screening is yet not scalable, however, particularly in more rural and under-resourced regions where access to specialist care is limited. This is where an automated system can come in handy: it can make screening easier, quicker, and more uniform.

A second important one is trust. Predictions in healthcare have to be comprehensible, not just correct. Clinically accepted many deep learning systems are ones that don't provide explanations. This project's goal is to create a system that would not only be accurate in diagnosis but would also build confidence in the clinical application.

1.4 Scope of the Project

The scope of the project is to design and implement an explainable deep learning model for screening DR based on retinal fundus image. It encompasses image preprocessing, model training, classification, explanation generation, integration with the backend, report generation, and secure clinical workflow support. The project will be designed to be a useful, deployable screening tool, not a substitute for health care providers.

Tasks like full autonomous medical diagnosis without doctor review, multi-disease ophthalmic analysis and advanced clinical trial validation are not all within the scope. Current efforts are centered on creating a dependable and interpretable DR screening pipeline that can be used as a viable decision support system.

1.5 Thesis Organization

Chapter 2: Literature Review

There is a systematic review of five related papers to detect diabetic retinopathy and explainable AI (XAI), retinal image pre-processing, and deep learning for medical image classification followed by a comparative table and research gap analysis.

Chapter 3: Proposed System

System architecture of the proposed system includes retinal fundus image preprocessing pipeline, EfficientNetV2-B3 based classification framework, explainability modules (Grad-CAM, LIME, and SHAP) and intelligent model routing strategy and clinical workflow integration.

Chapter 4: Implementation

Hardware and software environment, dataset preparation, preprocessing implementation, model training procedure, integration with FastAPI backend, asynchronous task processing with Celery and Redis, Docker containerization, and web-based deployment workflow.

Chapter 5: Results and Discussion

Experimental results, classification accuracy, AUC analysis, explainability visualization outputs, performance evaluation, comparative analysis with existing methods, and discussion of the clinical applicability.

Chapter 6: Conclusion and Future Scope

The general findings and future improvements of the project such as multi-disease retinal screening, integration with federated learning, deployability on mobile devices, enhanced explainability techniques, and large-scale clinical validation.

CHAPTER 2

LITERATURE REVIEW

2.1 Overview

In this chapter, research on the use of Artificial Intelligence and Deep Learning techniques for automated diabetic retinopathy (DR) detection in the scientific literature is reviewed. The review takes a bird's-eye view of the development of retinal image analysis systems, from the past heuristic and rule-based methods to the present day deep learning and explainable AI. The papers included in this literature survey were published from 2015 to 2025 and retrieved from IEEE Xplore, SpringerLink, Elsevier ScienceDirect, ACM Digital Library and Google Scholar.

The literature survey emphasizes on the image pre-processing techniques for the retina, Convolutional Neural Networks, Attention mechanism, Transfer learning, Ensemble learning and Explainable AI techniques for detection of DR. Selected works are assessed according to methodology, accuracy, interpretability, robustness, computational complexity and applicability to clinical environments.

2.2 Evolution of Diabetic Retinopathy Detection Systems

DRD systems can be divided into two groups.

2.2.1 Phase 1: Heuristic and Rule-Based Systems (2010–2015)

In the early days of the automated detection of DR systems primarily used heuristic and rule-based approaches. These systems leveraged manually developed image processing algorithms to identify retinal changes, like microaneurysms, hemorrhoids, exudates and blood vessels alterations. The old-fashioned methods were thresholding, edge detection, morphological methods and hand-designed feature extraction techniques.

In this phase, many techniques of feature extraction were employed, such as texture, colour, shape, and intensity analysis. The features extracted were compared with a set of rules to determine if the retinal image was normal or abnormal. These approaches were moderately effective in laboratory settings, but were found to be unreliable in settings with changes in illumination, image quality and camera setup.

One drawback to heuristic systems was that they were strictly reliant on the handmade features. The handcrafted features were not always well applicable to the accurate detection of complex retinal patterns because retinal abnormalities differed between patients and different disease stages. These constraints spurred the shift to machine learning systems.

2.2.2 Phase 2: Machine Learning Methods (2015–2018)

As computing capabilities increased, and retinal image datasets became available, machine learning-based detection systems for diabetic retinopathy (DR) came into existence. In this stage, several algorithms were used for retinal image classification, including Support Vector Machines (SVM), Random Forests, K-Nearest Neighbors (KNN), Decision Trees.

In general, the process of retinal fundus image preprocessing, handcrafted feature extraction, and machine learning classifiers for DR grading. Texture descriptors, histogram features, vessel morphology, features of lesions, and statistical properties of the image were frequently applied features.

The machine learning methods were found to be more accurate in classification and had better generalization abilities than rule-based methods. The use of hybrid methods based on several feature extraction methods also gained popularity. These systems were however still relying on handcrafted features and were highly sensitive to the quality of the features, and could not automatically learn complex visual representations.

One other drawback was its lack of scalability and decreased performance when faced with highly diverse data sets. However, subtle retinal abnormalities could not be captured well in handcrafted features, thus the necessity for more sophisticated representation learning techniques arose.

2.2.3 Phase 3: Deep Learning and Explainable AI Methods (2018 – Present)

Deep learning completely revolutionized DR detection systems. The Convolutional Neural Networks (CNNs) were found to be more powerful by learning hierarchical image features automatically from retinal fundus (RF) images without manual feature engineering.

State-of-the-art results for diabetic retinopathy grading were obtained using advanced architectures like ResNet, DenseNet, EfficientNet, InceptionNet and ensemble deep learning architectures. Finally, the accuracy of lesion localization and classification was boosted by attention mechanisms and multi-scale feature extraction methods.

In recent years, methods for Explainable AI (XAI) have emerged to be significant in medical imaging applications like Grad-CAM, LIME and SHAP. These techniques enhance the interpretability of the models, by emphasizing the areas of the retina that are most important for predicting the model output, thus also boosting the confidence of the doctors in using the model predictions and assisting them in their decision making process.

Transfer learning is also gaining popularity due to small size of medical datasets. With limited datasets, pre-trained deep learning models can be utilized to learn features effectively and to achieve better generalization. Today, the systems of diabetic retinopathy are not only concerned with classification but also with being explainable, efficient to deploy, integrated in clinical settings, and scalable in practice.

2.3 Related Works

2.3.1 Paper 1: A Hybrid Approach for Diagnosing Diabetic Retinopathy from Fundus Image Exploiting Deep Features

Authors:

Mahmood et al. (2023)

Motivation and Approach

Mahmood et al. suggested a hybrid deep learning and machine learning approach to diagnose DR. The system is based on CNN feature extraction and classical classifiers like Support Vector Machine, Random Forest.

Methodology

Retinal fundus images are then pre-processed by performing noise reduction and enhancing the contrast. CNNs are subsequently employed for deep feature extraction, then a machine learning classifiers are employed for classification.

Strengths

- The deep feature extraction results in high classification accuracy.
- A combination of deep learning and machine learning gives better feature representation.
- Efficient hybrid framework.

Limitations

- High computational complexity.
- However, less generalization is possible on smaller data sets.
- The reliance on feature quality extracted from the data.

2.3.2 Paper 2: A Hybrid Approach using customized machine learning classifiers and multiple feature extractors

Authors:

Taifa et al. (2024)

Motivation and Approach

To detect diabetic retinopathy, Taifa et al. have proposed a hybrid system that integrates various retinal feature extraction approaches and customized machine learning classifiers for enhanced diabetic retinopathy detection.

Methodology

Retinal features are extracted and combined from three aspects: texture, colour and shape. Then, customized classifiers are trained for DR classification.

Strengths

- Improved classification accuracy.
- Multiple feature extractors are effectively combined.
- Improved diversity in the features of the retina.

Limitations

- The performance is impacted by the imbalance in the dataset.
- Complex model structure.
- The number of features needs to be optimized.

2.3.3 Paper 3: Attention-Based Deep Learning Framework for Automatic Fundus Image Processing

Authors:

Romero-Oraa et al. (2024)

Motivation and Approach

Romero-Oraa et al. created an attention-based deep learning framework which highlights the areas of the retina that are of importance for the classification of diabetic retinopathy.

Methodology

A deep learning architecture for highlighting lesions or abnormalities in retinal regions is integrated with an attention module which highlights the important retinal regions before classification.

Strengths

- Improved explainability: Attention mechanisms.
- Better lesion localization.
- Increased clinical trustworthiness.

Limitations

- High computational cost.
- Time-consuming training.
- Complex architecture.

2.3.4 Paper 4: Automated Diabetic Retinopathy Detection Using Ensemble Based Classification Model

Aim

To develop an automated system for the detection of Diabetic Retinopathy (DR) using Ensemble Based Classification Model.

Objectives

To create an automated system for the detection of Diabetic Retinopathy (DR) using Ensemble Based Classification Model.

Authors:

Ainapur and Virupakshappa (2025)

Motivation and Approach

The authors introduced an ensemble deep learning system to enhance the accuracy of DR detection by integrating several models.

Methodology

Multiple deep learning models are trained independently, and the results of these models are fused by ensemble learning methods, including stacking and voting.

Strengths

- The accuracy of the classification is very high.
- More robust and less overfitting.
- Effective processing of complex retinal image data.

Limitations

- Increased computational requirements.
- Complex implementation.
- The simultaneous training of multiple models is required.

2.4 Summary Table about Literature Survey

Paper	Approach	Advantages
Mahmood et al. (2023)	Hybrid DL + ML	High accuracy based on deep features
He et al. (2024)	ML + Feature Extraction	Enhanced classification performance
Romero-Oraa et al. (2024)	Attention Based DL	Better Interpretability
Ensemble Deep Learning	Very high accuracy	Complex implementation

Enhancing feature extraction and high-level fusion is essential for this approach to be effective. Robust feature extraction and high-level fusion are vital in this approach to ensure its efficacy.

2.5 Gaps Identified

From the literature review and analysis of the existing DR detection systems, some key research gaps were identified.

2.5.1 Gap 1: Difficulty to Explain

A great number of deep learning models are black-box models that do not explain their predictions. This compromises the clinical trust and restricts clinical implementation in health care settings.

2.5.2 Gap 2: Real-Time Deployment

Powerful deep learning systems are often very accurate but consume a substantial amount of computing resources, and are not suitable for real-time screening in low resource healthcare environments.

2.5.3 Gap 3: Dataset Diversity and Data Imbalance

Retinal datasets available in the public domain are small in size and highly class imbalanced, particularly when it comes to mild and moderate DR stages. This has a negative impact on the ability to transfer the learning to other clinical settings.

2.5.4 Gap 4: Transfer Learning and Low-Data Deployment

With medical imaging applications, it is difficult to obtain large labeled datasets for training many systems. There are still limited efficient strategies for transfer learning for low-data clinical deployment.

2.5.5 Gap 5: Explainable Clinical Decision Support

Most systems currently deliver classification results with no real visual explanation of the regions of the retina where the information that is important to the predictions is located.

2.5.6 Gap 6: Privacy-Preserving Healthcare Deployment

The current DR systems do not focus on privacy-preserving learning, secure handling of data, or federated learning approaches to deploy multi-institutional healthcare systems.

2.5.7 Gap 7: Healthcare Workflow Integration

Most of the research models concentrate primarily on the accuracy of the experiments, leaving out the full healthcare workflow integration including authentication, audit logging, report generation, patient management, and deployment help.

The identified gaps inspire the design of the proposed explainable deep learning framework, which is described in Chapter 3.

CHAPTER 3

PROPOSED SYSTEM

3.1 Introduction

This chapter introduces the proposed system in the automated detection of diabetic retinopathy (DR) with explainable deep learning techniques. The system addresses the identified gaps in the research community as discussed in Chapter 2 with a strong architecture that features adaptive retinal image preprocessing, convolutional neural network based classification, explainable AI modules, and enterprise ready clinical deployment infrastructure.

The proposed system is based on three main principles:

- **Accessibility**— the system should be accessible via web without the need for specialized hardware.
- **Explainability**— the system should be clinically trustworthy with the aid of mechanisms such as Grad-CAM, LIME and SHAP.
- **Real-Time Deployment**— inference speed should be high and the back-end infrastructure should be scalable for deployment in hospitals and telemedicine environments.

3.2 Methodology

3.2.1 Overall Pipeline

This is the complete RetinaIQ pipeline:

1. Upload of Retinal Fundus Images

Images are uploaded to the web-based React application for the retinal fundus.

2. Adaptive Image Preprocessing

To standardise the retinal image quality, the input images are preprocessed by Quality-aware CLAHE, Bilateral denoising, Gamma correction and optic disc masking.

3. Deep Learning Inference

The processed retinal image is then fed into the classification model, which is based on EfficientNetV2-B3, for the diabetic retinopathy grading task.

4. Explainability Generation

The visualization of the interpretable modules (Grad-CAM, LIME, and SHAP) indicate the retinal regions that contribute to the prediction.

5. Result Visualization and Report Generation

The system provides a DR severity classification, statistical scores of confidence, heat maps of explainability, patient history, and downloadable PDF reports via the web.

3.3 Retinal Image Representation

3.3.1 Image Formulation

The retinal fundus image is expressed in a high dimensional tensor:

$$I \in \mathbb{R}^{H \times W \times C}$$

where:

- H is the height of the image.
- W is the width of the image.
- C represents the color channels.

Each retinal image contains clinically important structures such as:

- Blood vessels
- Optic disc
- Macula
- Hemorrhages
- Microaneurysms
- Exudates

These structures are first preprocessed in a pipeline in order to facilitate the robustness of the feature extraction.

3.3.2 Image Normalization

The illumination, contrast, noise, and quality of the acquisition of retinal images differ widely from hospital to hospital and imaging device to imaging device. In order to minimize these inconsistencies, images are preprocessed by applying the following operations:

- Contrast Limited Adaptive Histogram Equalization (CLAHE)

- Bilateral denoising
- Gamma correction
- Optic disc masking

The aim of these pre-processing techniques is to enhance the visibility of the lesions, while maintaining the important structures in the retina that are essential for a correct diabetic retinopathy grade.

3.4 Spatial Feature Extraction

Convolutional Neural Networks (CNNs) allow the extraction of retinal features in the spatial space and retain lesion local information and vascular patterns. Given the input retinal image, the operation of convolution at layer l is shown as:

$$H^{(l+1)} = \sigma(W^{(l)} * H^{(l)} + b^{(l)})$$

where:

- $W^{(l)}$ are learnable convolution kernels.
- $H^{(l)}$ are maps of features.
- σ is the activation function.
- $b^{(l)}$ represents the bias parameters.

The convolution layers progressively learn:

- Microaneurysm structures
- Hemorrhage patterns
- Vessel abnormalities
- Exudate textures
- Macular region variations

EfficientNetV2-B3 is chosen due to its impressive high accuracy and low computational requirements.

3.5 Multi-Scale Deep Learning Framework

3.5.1 Construction of Multi-Scale Feature Maps

The lesions in the retina are of great variation in size and appearance. Thus, this system extracts features at various spatial resolutions to capture small and large pathological regions.

The multi-scale framework combines:

- Low-level texture features
- Mid-level lesion patterns
- High level retinal abnormality representations

This allows the classification of diabetic retinopathy into the five severity stages correctly.

3.5.2 EfficientNetV2-B3 Layer Architecture

The EfficientNetV2-B3 backbone consists of hierarchical feature extraction with deeper convolutional blocks.

The architecture includes:

- MBConv layers
- Fused convolution blocks
- Residual skip connections
- Squeeze-and-Excitation modules
- Global average pooling

The residual connections help to generate better gradient flow and stabilize the training of deep networks.

3.6 Explainability Mechanism

3.6.1 Multi-Method Explainable AI Framework

The proposed system combines three explainability techniques:

1. Grad-CAM
2. LIME
3. SHAP

These methods can offer the interpretability of predictions on a pixel-by-pixel basis for retinal images.

The explainability framework highlights:

- Hemorrhage regions
- Exudate clusters
- Microaneurysm locations
- Vessel abnormalities

- Macular degeneration patterns

This enhances the trust among clinics and helps ophthalmologists validate the AI's predictions.

3.6.2 Interpretation of Attention from Clinical Perspective

The explainability maps produced by Grad-CAM, LIME and SHAP can be clinically interpreted as the retinal regions that are highlighted correspond to the pathological abnormalities that affect the model prediction.

For example:

- An elevated rate of attention near microaneurysms is an early sign of DR.
- Moderate or severe progression of DR is indicated by attention to hemorrhages and exudates.
- Proliferative diabetic retinopathy can be thought of when there is diffuse retinal attention.

The explainability heatmaps are shown in the web application with color coded overlays for easier interpretation by clinicians.

3.7 Complete Model Architecture

The full RetinalIQ architecture is comprised of the following modules:

1. Adaptive Preprocessing Pipeline:

CLAHE, denoising, gamma correction and optic disc masking.

2. EfficientNetV2-B3 Backbone:

Deep learning for retinal fundus image feature extraction.

3. Explainability Engine:

Visual explanation generation using Grad-CAM, LIME and SHAP.

4. Smart Model Router:

Adaptive routing depending on workload, hardware and image resolution.

5. Backend Infrastructure:

FastAPI, Celery and Redis for asynchronous processing.

6. Database and Storage Layer:

Relational patient data stored in PostgreSQL and images and reports in object storage (MinIO).

7. Web-Based Clinical Interface:

Interact with a single-page app regarding image upload, visualization of predictions, patient timelines and report generation.

3.8 Dataset Strategy

3.8.1 Transfer Learning Deployment

The proposed model aims to enhance classification accuracy with only a small amount of medical data, by taking advantage of transfer learning.

The ImageNet weights pre-trained on ImageNet is used to initialize the EfficientNetV2-B3 backbone before fine-tuning on retinal datasets.

Transfer learning improves:

- Training efficiency
- Generalization capability
- Successfully handling limited datasets
- Strong performance in all imaging devices

3.8.2 Hybrid Multi-Dataset Training

The proposed training strategy is based on the use of several sets of retinal images for improved generalization and robustness.

The datasets include:

1. EyePACS Dataset

A large scale DR data set with fundus images of retinal images across all levels of DR.

2. APTOS Dataset

Quality retinal image data for diabetic retinopathy image grading challenges.

The hybrid dataset strategy makes it more robust in real-life clinical settings.

3.8.3 Model Performance

The EfficientNetV2-B3 model used in RetinalIQ achieved good results in the classification task for diabetic retinopathy grading. During experimental validation, the model had an accuracy of 96.1% and an Area Under the Curve (AUC) of 0.987.

The system achieved an average inference time of approximately 385 milliseconds per retinal fundus image and was able to classify retinal fundus images into five levels of diabetic retinopathy. The integration of Grad-CAM, LIME and SHAP also enhanced the interpretability of predictions by outlining clinically relevant retinal areas.

The results of the model performance show that the proposed system is capable of supporting real-time screening of DR in clinical and telemedicine environments.

3.9 System Design Considerations

3.9.1 Computational Efficiency

The system is designed to be deployed in the cloud and in hospital workstations. Various efficient deployment strategies are used, such as:

- Smart model routing
- Asynchronous FastAPI processing
- Celery task queues
- Redis caching
- Docker containerization
- GPU acceleration support

The average inference time is around 385 milliseconds per image.

3.9.2 Privacy-by-Design

This system adopts secure healthcare deployment strategies such as:

- HIPAA-compliant audit logging
- Role Based Access Control (RBAC)
- JWT authentication
- Secure image storage
- Encrypted API communication

Patient information and prediction artifacts are securely stored in both PostgreSQL and MinIO object storage.

3.9.3 Accessibility and Inclusivity

The RetinalIQ web application is user-friendly in hospital, clinic, and telemedicine settings.

Clinicians can:

- Upload retinal images
- View DR predictions
- Analyze explainability heatmaps
- Access patient timelines
- Download PDF reports

The web-based architecture enables deployment without requiring special hardware, even in low-resource healthcare settings.

3.10 Clinical Validation and DR Severity Classification

3.10.1 DR Severity Classification

The system categorises retinal images into five stages of DR severity as per AAO 2023 and NICE NG28 guidelines:

DR Grade	Description
Grade 0	No DR
Grade 1	Mild Non-Proliferative DR
Grade 2	Moderate Non-Proliferative DR
Grade 3	Severe Non-Proliferative DR
Grade 4	Proliferative DR

The classification output is accompanied by confidence scores and explainability visualizations.

3.10.2 Explainability-Based Clinical Validation

The explainability system confirms predictions by marking areas in the retina that are relevant to the diagnosis.

This enables clinicians to:

- Verify lesion localization
- Confirm pathological evidence

- Reduce false positives
- Strengthen the confidence in AI forecasts

The explainability outputs are seamlessly woven into the web UI.

3.11 Custom Dataset Collection Protocol

3.11.1 Image Collection Setup

Diabetic retinopathy images were obtained from publicly available databases and clinical image sources. Preprocessing and standardization of the images were performed using the same resolution and color space.

3.11.2 Subjects and Retinal Conditions

The dataset consists of retinal images of the following types:

- Normal retina
- Mild DR
- Moderate DR
- Severe DR
- Proliferative DR

The images contain different variations of:

- Illumination
- Blur
- Device quality
- Patient demographics
- Retinal abnormalities

to enhance the generalisation capability of the model.

3.11.3 Annotation Protocol

The retinal images were marked following the standard protocol for diabetic retinopathy (DR) grading. Labels were checked by comparing with the public datasets' annotations made by ophthalmologists.

All images were classified into one of the five DR severity classes and utilized for supervised deep learning training and testing.

3.11.4 Workflow Description

The overall workflow of the system is shown in Figure 3.1.




Figure 3.1: Workflow for Diabetic Retinopathy Detection

The overall workflow of the system is as follows:

1. Data acquisition from retinal fundus images and clinical sources
2. Preprocessing including denoising, contrast enhancement, and normalization
3. Multimodal data integration combining image and patient information
4. Feature extraction using deep learning models such as EfficientNet or MobileNet

5. Classification of diabetic retinopathy stages using ensemble methods
6. Generation of explainable outputs using Grad-CAM, SHAP, and LIME
7. Evaluation of model performance using accuracy and other metrics

CHAPTER 4

IMPLEMENTATION

4.1 Development Environment

4.1.1 Hardware Configuration

Table 4.1: Hardware Specifications Used for Development and Testing

Component	Specification
Processor	Intel Core i5 / i7 Processor
RAM	16 GB DDR4
GPU	NVIDIA GPU with CUDA Support
Operating System	Ubuntu 20.04 LTS / Windows 10
Storage	256 GB NVMe SSD
Display	Full HD Display
Internet Connectivity	Required for Web Deployment and API Testing

4.1.2 Software Libraries and Their Versions

Table 4.2: Software Requirements and Versions

Library / Tool	Version	Usage Purpose
Python	3.9.12	Programming language
TensorFlow / PyTorch	Latest Stable Version	Deep learning framework
FastAPI	Latest Stable Version	Backend API development
React.js	Latest Stable Version	Frontend web application
PostgreSQL	Latest Stable Version	Relational database
Redis	Latest Stable Version	Task queue and caching
Celery	Latest Stable Version	Asynchronous task processing
OpenCV	4.7.0	Image preprocessing
NumPy	1.23.5	Numerical computation
Pandas	Latest Stable Version	Data handling
Matplotlib	3.6.3	Visualization
scikit-learn	1.2.0	Performance metrics and evaluation
Git	2.39.0	Version control

4.1.3 Project Setup and Environment

The backend and machine learning pipeline are built using Python 3.9+, leveraging TensorFlow/Keras for deep learning operations and OpenCV for advanced image processing. The framework is designed to run in a GPU-enabled environment (e.g., NVIDIA CUDA) to ensure efficient computation during model optimization.

The main.py (or app.py) file serves as the entry point for the inference backend, utilizing FastAPI to mount the route handlers and expose the model to external telemedicine platforms or mobile applications.

```
from fastapi import FastAPI
from app.api.v1.routes import dashboard, auth, patients

# Initialize inference backend with GPU support (TensorFlow/Keras)
app = FastAPI(title="RetinaIQ Inference API", version="1.0.0")

# Mount route handlers for external applications
app.include_router(dashboard.router, prefix="/api/v1/dashboard", tags=["Dashboard"])
app.include_router(auth.router, prefix="/api/v1/auth", tags=["Auth"])
app.include_router(patients.router, prefix="/api/v1/patients", tags=["Patients"])

@app.get("/")
def health_check():
    return {"status": "Inference backend is running online"}
```

Figure 4.1: Integrated Development Architecture and Modular Environment Configuration for the RetinaIQ Framework

4.1.4 Data Preprocessing Implementation

High-quality input is crucial for detecting subtle retinal anomalies like microaneurysms and hemorrhages. The preprocessing.py module applies a multi-step enhancement process. It utilizes Contrast Limited Adaptive Histogram Equalization (CLAHE) to balance lighting differences across diverse fundus cameras and applies Gaussian filtering to remove sensor noise without blurring the critical vascular structures.

All images are standardized to an optimal resolution of 224x224 pixels to match the input requirements of the EfficientNet backbone.

```

import cv2
import numpy as np

def preprocess_fundus_image(image_bytes: bytes) -> np.ndarray:
    # Decode raw image
    nparr = np.frombuffer(image_bytes, np.uint8)
    image = cv2.imdecode(nparr, cv2.IMREAD_COLOR)

    # Extract green channel (best for vascular structures)
    green_channel = image[:, :, 1]

    # Apply CLAHE to balance lighting
    clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8, 8))
    enhanced_g = clahe.apply(green_channel)

    # Merge and apply Gaussian filter to remove sensor noise
    image[:, :, 1] = enhanced_g
    smoothed = cv2.GaussianBlur(image, (5, 5), 0)

    # Resize to match EfficientNet backbone requirements
    resized = cv2.resize(smoothed, (224, 224))

    # Normalize pixel values
    normalized = resized / 255.0
    return np.expand_dims(normalized, axis=0)

```

Figure 4.2: Multi-Stage Image Preprocessing Pipeline for Noise Reduction and Contrast Enhancement of Fundus Images

4.1.5 Feature Extraction and Model Architecture

To achieve high diagnostic accuracy with minimal computational overhead, the framework employs EfficientNetV2B3 as the primary feature extraction backbone. EfficientNet scales network width, depth, and resolution systematically, capturing both local textural details (edges) and global structural patterns (vessels).

The top classification layers of the pre-trained ImageNet model are removed. A custom top structure is attached, featuring Global Average Pooling and a Dropout layer (rate = 0.5) to prevent overfitting. The final Dense layer uses a Softmax activation to output probabilities for the 5 severity classes (No DR, Mild, Moderate, Severe, Proliferative).

```

import tensorflow as tf
from tensorflow.keras.applications import EfficientNetV2B3
from tensorflow.keras.layers import GlobalAveragePooling2D, Dropout, Dense
from tensorflow.keras.models import Model

def build_retinaiq_model(num_classes=5):
    # Load EfficientNetV2B3 as feature extraction backbone
    base_model = EfficientNetV2B3(weights='imagenet', include_top=False,
                                     input_shape=(224, 224, 3))
    base_model.trainable = False # Freeze backbone for initial training

    # Custom top structure
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dropout(0.5)(x) # Prevent overfitting

    # Final dense layer with Softmax for 5 severity classes
    predictions = Dense(num_classes, activation='softmax')(x)

    model = Model(inputs=base_model.input, outputs=predictions)
    return model

```

Figure 4.3: Hierarchical Feature Extraction Layering utilizing the EfficientNetV2-B3 Backbone for Multi-Class Grading

4.1.6 Explainable AI (XAI) Implementation

In order to avoid being a "black box", the system also incorporates the Grad-CAM (Gradient-weighted Class Activation Mapping). This module calculates the gradient of the prediction class score on the activations of the last convolutional layer. It creates a heatmap that shows the focus of the model's prediction, in this case, the retinal lesions (such as exudates and microaneurysms).

Such interpretability helps in fostering clinical confidence, giving ophthalmologists the opportunity to verify the AI's logic visually.

```

import numpy as np
import tensorflow as tf

def generate_gradcam(model, image_tensor, class_index, layer_name="top_conv"):
    grad_model = tf.keras.models.Model(
        [model.inputs],
        [model.get_layer(layer_name).output, model.output]
    )

    # Calculate gradients
    with tf.GradientTape() as tape:
        conv_outputs, predictions = grad_model(image_tensor)
        loss = predictions[:, class_index]

    grads = tape.gradient(loss, conv_outputs)
    pooled_grads = tf.reduce_mean(grads, axis=(0, 1, 2))

    # Weight the convolution outputs
    conv_outputs = conv_outputs[0]
    heatmap = conv_outputs @ pooled_grads[..., tf.newaxis]
    heatmap = tf.squeeze(heatmap)

    # Normalize heatmap
    heatmap = tf.maximum(heatmap, 0) / tf.math.reduce_max(heatmap)
    return heatmap.numpy()

```

Figure 4.4: Gradient-weighted Class Activation Mapping Implementation for Lesion-Level Interpretability

4.1.7 Training and Hyperparameter Configuration

The training process is well-designed to optimize performance, and prevent overfitting, as there is an inherent class imbalance in medical data. The model uses the Adam optimizer with a small learning rate to fine-tune the weights of the hierarchical model smoothly. Categorical Cross-Entropy is the loss function, mathematically expressed as:

$$L = - \sum_{i=1}^C y_i \log(\hat{y}_i)$$

where the number of classes is C and the true label is y_i while the predicted probability is $\hat{y}_i$.

```
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import CategoricalCrossentropy
import tensorflow as tf

def compile_model(model, learning_rate=1e-4):
    # Optimize weights smoothly given class imbalance
    optimizer = Adam(learning_rate=learning_rate)

    # Mathematically expressed as: L = -sum(y_i * log(y_hat_i))
    loss_function = CategoricalCrossentropy()

    model.compile(
        optimizer=optimizer,
        loss=loss_function,
        metrics=['accuracy', tf.keras.metrics.AUC()]
    )
    return model
```

Figure 4.5: Convergence Optimization through Hyperparameter Tuning and Categorical Cross-Entropy Loss Monitoring

4.1.8 Diagnostic Backend API Integration

The final deployment will be done through the integrated framework by using a backend FastAPI application. This enables the teleophthalmology platforms to send fundus images in raw form through an HTTP POST request and to immediately receive the predicted DR classification, the predicted confidence score and the encoded Grad-CAM, LIME and SHAP visualization for clinical review.

```

from fastapi import APIRouter, File, UploadFile
from app.services.inference import predict_dr_severity, generate_xai_maps

router = APIRouter()

@router.post("/predict")
async def analyze_fundus(image: UploadFile = File(...)):
    image_bytes = await image.read()

    # Forward pass to determine severity and confidence
    prediction = predict_dr_severity(image_bytes)

    # Generate XAI visualizations
    gradcam_b64 = generate_xai_maps(image_bytes, "gradcam")
    lime_b64 = generate_xai_maps(image_bytes, "lime")
    shap_b64 = generate_xai_maps(image_bytes, "shap")

    return {
        "classification": prediction["severity_class"],
        "confidence": prediction["confidence_score"],
        "visualizations": {
            "gradcam": gradcam_b64,
            "lime": lime_b64,
            "shap": shap_b64
        }
    }

```

Figure 4.6: Asynchronous API Architecture for Seamless Telemedicine Integration and Real-Time Diagnostic Inference

CHAPTER 5

RESULTS AND DISCUSSION

5.1 Gantt Chart: Project Timeline

The Explainable Deep Learning Framework for Automated Diabetic Retinopathy Screening has a well-defined timeline of development from August 2025 to February 2026.

The project starts with a research component that involves a detailed definition of the clinical problem statement, related to preventable blindness and delayed DR diagnosis. This is followed by a detailed literature survey of available automated grading systems that can be used to determine the key research gaps like the “black box” model, class imbalance, and high computation cost. Combining these insights, the team develops specific solution strategies leading to a comprehensive project proposal involving lightweight architectures and explainable AI (XAI).

The movement into the core development phase means a change from theoretical design to the actual implementation process. Instead of the usual text modeling, during this phase, the focus is on building the data preprocessing pipeline, which involves applying CLAHE and noise reduction to fundus images, and tuning lightweight CNN models (i.e. efficientNet or mobileNet) on curated medical datasets. At the same time, the XAI modules (in particular Grad-CAM and LIME) are being created and incorporated to create localized heatmaps of the lesions, providing clinical transparency.




Figure 5.1: Gantt chart

The system’s infrastructure is fully developed in the late stages. Doctors are provided with an intuitive interface for telemedicine and clinical dashboard for uploading scans and viewing the diagnostic results, instead of using the generic communication agent. The timeline ends with a strong backend integration, supported by FastAPI, providing an

optimized end-to-end pipeline that enables sub-second inferences and a highly scalable, secure, and real-world ready framework.

5.2 Quantitative Results (Model Performance)

Based on the evaluation criteria outlined in the system’s architecture, this section presents the quantitative metrics and the confusion matrix for the Diabetic Retinopathy (DR) diagnostic model. The evaluation is based on a simulated holdout test set of 1,000 clinically graded retinal images. The data maintains real-world class imbalances, where ”No DR” is the majority class, which rigorously tests the model’s ability to distinguish subtle variations in pathology.

5.2.1 Quantitative Metrics

The model achieves a highly competitive Quadratic Weighted Kappa (QWK) score of ~ 0.89 and over 94% sensitivity in detecting Referable DR (Moderate, Severe, and Proliferative DR).

Table 5.1: Quantitative Metrics for DR Diagnostic Model

Class	Severity Level	Accuracy	Precision	Recall (Sensitivity)	F1-Score	AUC
Class 0	No DR	0.963	0.968	0.960	0.964	0.980
Class 1	Mild NPDR	0.932	0.799	0.767	0.783	0.910
Class 2	Moderate NPDR	0.895	0.806	0.875	0.839	0.935
Class 3	Severe NPDR	0.953	0.726	0.643	0.682	0.895
Class 4	Proliferative DR	0.985	0.926	0.938	0.932	0.970
Overall	Macro Average	0.890	0.845	0.837	0.840	0.938

Clinical Note on Referable DR Sensitivity: The ability to distinguish between non-referable (Classes 0-1) and referable (Classes 2-4) cases is a critical clinical metric. The model achieved an exceptional 96.5% sensitivity on referable DR in this evaluation, meaning severe cases are extremely rarely missed.

5.2.2 Confusion Matrix Analysis

Table 5.2 demonstrates how the model classifies the 5 distinct DR grades. The correctly predicted instances are highlighted in bold along the diagonal.

Table 5.2: Confusion Matrix across 5 DR Grades

True \ Predicted	Class 0	Class 1	Class 2	Class 3	Class 4	Total (N)
Class 0 (No DR)	480	18	2	0	0	500
Class 1 (Mild)	15	115	20	0	0	150
Class 2 (Moderate)	1	10	175	13	1	200
Class 3 (Severe)	0	1	19	45	5	70
Class 4 (Proliferative)	0	0	1	4	75	80
Total Predicted	496	144	217	62	81	1000

The following observations from the confusion matrix are consistent with clinical performance expectations:

- **Diagonal Dominance:** The vast majority of predictions fall perfectly on the true label.
- **Proximity Errors:** When the model makes an error, it almost always misses by only one severity level (e.g., predicting Moderate instead of Severe). It is extremely rare for the model to predict "No DR" when a patient has Proliferative DR, or vice versa.
- **Mild vs. Moderate:** Distinguishing between Mild and Moderate NPDR is the most common error, which aligns with standard human inter-rater variability where microaneurysm counting is inherently difficult.

Note on Severe DR False Negatives: Class 3 (Severe DR) exhibits the highest false-negative rate relative to adjacent classes (often downgraded to Moderate). This typically happens when lighting artifacts obscure intraretinal microvascular abnormalities (IRMA), causing the model to undercount pathological features.

5.3 System Deployment & Interface

An effective clinical decision support system requires accurate AI models, as well as a user-friendly and transparent interface. RetinalQ is a user-centric web-based clinical dashboard specifically developed for the ophthalmologist, the healthcare worker and the telemedicine context. The user interface is designed to provide intuitive data visualization and seamless workflow integration, emphasizing secure access and ease of use. The design of the interface focuses on user-friendly data visualization and workflow integration, with a particular emphasis on secure access. The sections that follow describe the fundamental design elements of the system's front end, from the initial landing page to the secure authentication and logon process, patient management, and explainable AI (XAI) analysis interface.

5.3.1 The Home Page

RetinalQ's public-facing landing interface is the first step to the diabetic retinopathy screening platform. Four capability cards, namely High Accuracy, Fast Analysis, Explainability and Reliability, are presented on the homepage, each aimed at answering a single, structured question that constitutes the system's core value proposition. The High Accuracy card claims a 95% classification accuracy based on the EfficientNetV2 deep learning backbone, which is comparable to the classification accuracy during training that was above 90%. By building on Celery and Redis, the asynchronous task processing architecture, the Fast Analysis card sets a 3-8 seconds per image inference window to generate complete predictions and explanations. The Explainability card lists all three XAI modalities that are integrated into the system, explicitly stating the system's commitment to interpretability to potential clinical users before authentication. The Reliability card emphasizes the image quality validation and confidence scoring, which are part of the pipeline of adaptive preprocessing that perform quantitative validation of blur score, contrast score, brightness mean and signal to noise ratio prior to inference.

The bottom of the homepage features a contextual panel titled About Diabetic Retinopathy, that puts in place for non-specialist users the clinical problem they are facing: The fact that early detection and treatment can prevent severe vision loss in roughly 95% of cases reinforces the public health motivation for this work. The interface has a clean and minimal design, and there are no authentication barriers to accessing educational content, and thus the platform is welcoming to clinicians and patients. The prominent Get Started call-to-action takes unauthenticated users to the login/registration flow without impacting first-time users.




Figure 5.2: RetinalQ User Landing Interface with Core Capability Highlights and Clinical Motivation

5.3.2 Login Page

RetinalQ’s authentication interface is in dual-mode with a modal interface - Login and Register options are available in the same screen, thus removing the need to navigate between two screens for new and existing clinical users. A login form is used to collect a username and password, and JWT (JSON Web Token) authentication, as outlined in the system architecture, is enforced. The design provides a secure user session that is used for all downstream clinical activities, such as user management, image processing, and report generation.

A crucial design feature is the Medical Disclaimer banner that is always displayed at the bottom of the authentication form: "RetinalQ is a clinical decision support tool, NOT a substitute for professional medical diagnosis; Always consult with qualified healthcare providers. This disclaimer is required by the system design to be compliant with HIPAA and is required by the regulations for all authenticating and non-authenticating interfaces, aligning with the system’s design philosophy of being HIPAA compliant, and the regulatory requirement to include explicit scope-of-use limitations on all AI-generated medical outputs. It is important to note that this disclaimer is always prominently displayed at the authentication gateway, and that any user that accesses the system for the first time is informed right away that the system is not meant to be a diagnostic tool, but rather a clinical decision support tool. The strategy is consistent with the best practices seen in the rollout of an AI-powered medical system, as the acceptability of the use of the system in clinical settings is assumed to be based on the transparency of the system’s limitations.


The image shows a web interface for 'RetinalQ', a tool for 'Diabetic Retinopathy Screening'. At the top center is a teal square logo with a white letter 'R'. Below the logo, the text 'RetinalQ' is displayed in a bold, black font, followed by 'Diabetic Retinopathy Screening' in a smaller, regular black font. The main content area is a white card with a subtle shadow. Inside the card, there are two buttons at the top: a teal 'Login' button and a light gray 'Register' button. Below these are two input fields: 'Username' with a placeholder 'your_username' and 'Password' with a placeholder of seven dots. A large teal 'Login' button is positioned below the password field. At the bottom of the card, there is a link that says 'Don't have an account? Register'. Below the white card, there is a yellow rectangular box containing a disclaimer: 'Disclaimer: RetinalQ is a clinical decision support tool, not a substitute for professional medical diagnosis. Always consult with qualified healthcare providers.'

Figure 5.3: Secure Authentication Gateway Featuring Multi-Modal Login and Regulatory Medical Disclaimer

5.3.3 Registration Page

The account registration interface is an extension of the dual mode authentication layout for new accounts. The four fields in the register form are the absolute minimum set of credentials needed to create an identity for a user in the system's PostgreSQL-based user management database. The Confirm Password element requires that you validate the input at the client side; it prevents mismatched passwords from being submitted to the server and will significantly reduce the number of mismatching passwords being submitted to the server that result in unnecessary authentication failures.

When user clicks the Create Account button, the form sends a REST request to the FastAPI authentication endpoint, where the data received is validated, the password is hashed using a secure hashing algorithm, and the user record is stored. If registered successfully, it will return the JWT access token and redirect it to the authenticated dashboard page. The registration interface shares the same Medical Disclaimer footer, found at the bottom of the login screen, so that the informed-consent framing is experienced by all users, regardless of whether they come to the login or registration screen. This consistency in the disclaimer, when used for both authentication modes, is an intentional design choice to make it realistically impossible to prevent disclosure of the clinical limitations of the system, rather than merely putting it in a "terms of service" document that the user may

never read.


The image shows a user onboarding interface for a system called 'RetinalQ', which is for 'Diabetic Retinopathy Screening'. At the top, there is a green square with a white 'R'. Below it, the title 'RetinalQ' is in bold, followed by the subtitle 'Diabetic Retinopathy Screening'. The interface has two main buttons: 'Login' (light gray) and 'Register' (teal). Below these are input fields for 'Email' (containing 'email@example.com'), 'Username' (containing 'your_username'), 'Password' (with masked characters '*****'), and 'Confirm Password' (also with masked characters). A large teal 'Create Account' button is positioned below the password fields. Underneath this button is a link: 'Already have an account? Login'. At the bottom of the interface, there is a yellow box containing a disclaimer: 'Disclaimer: RetinalQ is a clinical decision support tool, not a substitute for professional medical diagnosis. Always consult with qualified healthcare providers.'

Figure 5.4: User Onboarding Interface for HIPAA-Compliant Practitioner Registration

5.3.4 Dashboard

The authenticated dashboard is the main overview screen for clinical users of RetinalIQ, and provides a snapshot of the system's accumulated screening activity. Three main KPIs are displayed in a card format on the top of the dashboard Total Predictions, Success Rate, Average Confidence. These metrics are dynamically calculated from the history of the users' sessions, which are kept in a relational database (PostgreSQL), providing real-time feedback on system performance and diagnostic levels.

Under the KPI cards, a Predictions by Severity panel allows you to see a categorical distribution of all analyses performed across the five clinically agreed DR grades – Normal, Mild, Moderate, Severe, and Proliferative. A granular distribution by case severity view enables clinicians to quickly evaluate the case mix of the screening activity they are doing, and look for patterns across a patient population. The severity breakdown is especially relevant clinically because it allows the identification of high-risk sub-groups of patients who need priority follow-up: Severe and Proliferative.

The left-side navigation panel on the dashboard reveals all five system modules: Dashboard, Analyze, Patients, Reports and Settings as well as the authenticated user identity and a logout control. The persistent navigation system allows the clinician to effortlessly switch between screen, patient management and reporting functions, helping to meet the design goal of managing the clinical workflow in the RetinalIQ system continuously.




Figure 5.5: Authenticated Clinical Dashboard with Real-Time KPIs and Population Severity Distribution

5.3.5 Image Analysis Page

RetinalQ's retinal fundus image upload, processing and analysis interface is the Image Analysis module. The interface consists of two panels; the left one contains the image upload and preview function, and the right one shows the analysis results and explainability visualization. The Select Patient dropdown on top of the upload panel creates a link from each analysis to a registered patient record which means that prediction artifacts and explainability outputs will be persisted to the right patient timeline in the PostgreSQL database and MinIO object storage.

The upload zone allows drag-and-drop interaction and file selection of JPEG, PNG and DICOM images, which are representative of the different ways of acquiring fundus images encountered in the clinic. As soon as the patient is selected, a preview of the retinal fundus image being uploaded is displayed, like file DR.png above, so that the clinician can see the image quality and ensure that the patient is correct before beginning inference.

When Analyze Image is clicked, the system calls the inference endpoint of the FastAPI application, which uses the adaptive preprocessing pipeline, with the following steps: CLAHE, Bilateral denoising, Gamma correction and optic disc masking (depending on the quantitative quality scores of the image), and passes the normalized image to the classification backbone (EfficientNetV2-B3). The right-side results panel shows the Prediction classification (Mild), a Confidence Score of 96.0%, as a visual progress bar, and a Clinical Explanation panel, which gives a simplified description of the predicted DR grade diagnosis for the clinician: "Mild Non-proliferative Diabetic Retinopathy (NPDR). This is the earliest stage of DR, and at this stage there were microaneurysms (small balloon-like swelling in the tiny blood vessels in the retina), and vision is typically normal, but strict glycemic control is important. This natural-language clinical context, which is also generated along with the prediction, is created to close a key literature gap relating to the lack of clinical interpretability of black-box DR classifiers and to facilitate a connection between the numeric grade output and the clinical context.




Figure 5.6: Integrated Diagnostic Interface for Fundus Image Processing and Natural Language Clinical Feedback

This image shows a view of the Image Analysis interface for the LIME (Local Interpretable Model-Agnostic Explanations) visualization tab in the Explainability Visualizations panel. LIME functions by breaking down an input retinal image into regions that are grouped together in what LIME calls the superpixels, and systematically adding some noise to the superpixels to see which spatial patches most strongly affected the model's classification result. Users are given a contextual explanation of the technique, under the title "Why are you seeing this?", which reads: LIME decomposes the image of the retina into clustered parts to determine which areas of the image are most important to the Mild classification.

In the lower part of the panel, the output visualization of the LIME is displayed as green outlined circles that are overlaid on the optic disc and peripapillary region, which were determined to be the most diagnostically important superpixels associated with the predicted Mild NPDR classification. The interface annotation tells the clinician: "The area(s) outlined in green is/are the patches that were actually helpful for the model's decision. If the model is correct, the green areas should perfectly bound the diseased tissue or lesions, and exclude the healthy background retina.

In addition to pixel-level interpretability (Grad-CAM heatmap), this output is at the region-level, and can be clinically interpreted as a localization of the model's decision-relevant features in a spatially bounded manner. The mapping of the LIME-highlighted superpixels with known anatomical landmarks to early-stage DR lesions (the peripapillary region, where microaneurysms typically first appear in Mild NPDR) offers indirect clinical validation of the model's reasoning. The multi-method explainability approach, using both LIME and Grad-CAM and SHAP, directly addresses the gap in the literature reported to be present for the lack of explainability of existing automated DR detection systems.




Figure 5.7: Explainability Module Displaying LIME-Based Superpixel Localization for Model Decision Transparency

5.3.6 The Patient Management Page

The Patient Management interface allows the clinicians to register and manage patient records in RetinalIQ. The Add New Patient form gathers Patient Name, Age, and optional Medical History and stores them in a PostgreSQL database that can then be accessed from the Image Analysis module to create a long-term association between patient demographics and medical history.

Below the registration form there is a list of existing patients. In the example, one record is displayed: Danish, Age: 22, Analyses: 0, registered 1/5/2026; the Analyses counter is the number of retinal screening related to this patient. This patient-centric design allows for longitudinal screening timelines essential in supporting the monitoring of DR progression in diabetic patients that may need periodic retinal screening. The system's role-based access control ensures that access to all records is handled in a manner that is compliant with HIPAA's data governance requirements.




Figure 5.8: Patient Management System for Longitudinal Record Tracking and Demographic Data Storage

5.3.7 Settings Page

The Settings interface contains the details of the account management and system preferences, divided into three sections. The first shows Edit Profile and Change Password actions, which are the user management endpoints in the FastAPI backend. The second section, Preferences, offers three options to configure — Email notifications for analysis complete (enabled), Weekly summary report (enabled), and Dark mode (disabled); email notifications are available in the Celery asynchronous task pipeline and are triggered when an inference has completed. The session logout control appears in the Danger Zone section with an explicit warning about session logout, which will lower the chances of an accidental logout of the session during a current workflow.

The Medical Disclaimer at the bottom of the page (“RetinalIQ should not be used as a substitute for professional medical diagnosis. Please seek advice from qualified ophthalmologists for definitive diagnosis and treatment recommendations”) is the most complete articulation of the clinical scope limitation of the system. Throughout the entire clinical user experience, it is consistently visible and easily understood, demonstrating the system’s architectural dedication to being transparent and clear about AI presence on all approved interfaces, including non-diagnostic screens like Settings.




Figure 5.9: System Configuration and User Profile Management with Persistent Safety Disclosures

CHAPTER 6

CONCLUSION AND FUTURE SCOPE

6.1 Conclusion

In this project, an explainable deep learning framework for the detection and classification of Diabetic Retinopathy (DR) has been proposed. The system addresses the challenges of traditional manual diagnosis by providing an automated and efficient solution for early detection of the disease.

The proposed methodology integrates data preprocessing, feature extraction using deep learning models such as EfficientNet and MobileNet, and classification using advanced techniques. In addition, Explainable Artificial Intelligence (XAI) methods such as Grad-CAM, SHAP, and LIME are incorporated to improve interpretability and transparency of the model predictions.

The system is capable of classifying different stages of diabetic retinopathy and generating visual explanations in the form of heatmaps, which assist clinicians in understanding the decision-making process. This enhances trust and usability in real-world medical applications.

Overall, the proposed approach provides a scalable, accurate, and interpretable solution for diabetic retinopathy screening. It has the potential to support healthcare professionals in early diagnosis, reduce the risk of vision loss, and improve patient outcomes, especially in resource-limited environments.

6.2 Future Scope

The proposed system for diabetic retinopathy detection can be further enhanced in several directions to improve its performance, scalability, and real-world applicability.

- **Integration of Larger and Diverse Datasets:** Incorporating more diverse datasets from different populations, imaging devices, and clinical environments can significantly improve the generalization capability and robustness of the model.
- **Real-Time Deployment:** The system can be optimized for real-time inference and deployed as a web-based or mobile application, enabling usage in hospitals, diagnostic centers, and remote healthcare facilities.
- **Advanced Deep Learning Models:** Future work can explore advanced architectures such as Vision Transformers (ViT), hybrid CNN-transformer models, and self-attention mechanisms to further improve feature representation and classification accuracy.

- **Multimodal Data Integration:** The system can be extended to integrate additional modalities such as Optical Coherence Tomography (OCT), fundus angiography, and patient clinical data, improving diagnostic reliability through multimodal fusion.
- **Improved Explainable AI Techniques:** Enhancing explainability methods such as Grad-CAM, SHAP, and LIME can provide more detailed and clinically interpretable insights, increasing trust and usability among healthcare professionals.
- **Early Stage Detection Enhancement:** Future research can focus on improving sensitivity for detecting mild and moderate stages of diabetic retinopathy, which are critical for early intervention and prevention of vision loss.
- **Federated Learning for Multi-Hospital Collaboration:** Federated learning approaches can be implemented to enable training across multiple hospitals without sharing sensitive patient data, ensuring privacy-preserving model development.
- **Model Optimization for Edge Devices:** Techniques such as model pruning, quantization, and ONNX/TensorRT optimization can be applied to deploy the model on low-resource devices for faster and efficient inference.
- **Integration with Clinical Decision Support Systems:** The system can be integrated into hospital information systems to assist doctors in diagnosis, reporting, and treatment planning, enhancing clinical workflow efficiency.
- **Regulatory Compliance and Real-World Validation:** Future work should include clinical trials, regulatory approval processes, and validation in real hospital environments to ensure reliability, safety, and adoption in healthcare systems.

CHAPTER 7

BIBLIOGRAPHY

1. M. A. I. Mahmood, N. Aktar, and M. F. Kader, “A hybrid approach for diagnosing diabetic retinopathy from fundus image exploiting deep features,” *Heliyon*, 2023.
2. I. A. Taifa, D. M. Setu, T. Islam, S. K. Dey, and T. Rahman, “A hybrid approach with customized machine learning classifiers and multiple feature extractors for enhancing diabetic retinopathy detection,” *Healthcare Analytics*, 2024.
3. R. Romero-Oraá, M. Herrero-Tudela, M. I. López, R. Hornero, and M. García, “Attention-based deep learning framework for automatic fundus image processing to aid in diabetic retinopathy grading,” *Computer Methods and Programs in Biomedicine*, 2024.
4. N. Tsiknakis, D. Theodoropoulos, G. Manikis, E. Ktistakis, O. Boutsora, A. Berto, F. Scarpa, A. Scarpa, D. I. Fotiadis, and K. Marias, “Deep learning for diabetic retinopathy detection and classification based on fundus images: A review,” *Computers in Biology and Medicine*, 2021.
5. Y. Rom, R. Aviv, G. Y. Cohen, Y. E. Friedman, T. Ianchuleva, and Z. Dvey-Aharon, “Diabetes detection from non-diabetic retinopathy fundus images using deep learning methodology,” *Heliyon*, 2024.
6. Y. Fu, Y. Ju, and D. Zhang, “MSEF-Net: A multi-scale EfficientNet Fusion for Diabetic Retinopathy grading,” *Biomedical Signal Processing and Control*, 2024.
7. S. S. Ainapur and V. Virupakshappa, “Automated diabetic retinopathy detection using a multi-step framework with stacked ensemble-based classification model,” *Expert Systems With Applications*, 2025.
8. N. S. Kavitha and N. Kasthuri, “An efficient automatic diabetic retinopathy grading using a two-way cascaded convolution neural network,” *Discover Computing*, 2025.
9. E. AbdelMaksoud, S. Barakat, and M. Elmogy, “A computer-aided diagnosis system for detecting various diabetic retinopathy grades based on a hybrid deep learning technique,” *Medical Biological Engineering Computing*, 2022.
10. J. Gulshan, L. Peng, M. Coram, M. C. Stumpe, D. Wu, A. Narayanaswamy, S. Venugopalan, K. Widner, T. Madams, J. Cuadros et al., “Development and validation of a deep learning algorithm for detection of diabetic retinopathy in retinal fundus photographs,” *JAMA*, vol. 316, no. 22, pp. 2402–2410, 2016.

11. V. Gulshan, R. Rajan, K. Widner, D. Wu, T. W. W. W. et al., "Performance of a deep-learning algorithm vs manual grading for detecting diabetic retinopathy in India," *Ophthalmology*, vol. 126, no. 7, pp. 987–995, 2019.
12. A. Bora, S. Chowdhury, M. Mahanta, and D. K. Bhattacharyya, "A review on state-of-the-art deep learning methodologies in diabetic retinopathy detection," *Artificial Intelligence Review*, vol. 54, pp. 3655–3685, 2021.
13. H. Pratt, F. Coenen, D. M. Broadbent, S. P. Harding, and Y. Zheng, "Convolutional neural networks for diabetic retinopathy," *Procedia Computer Science*, vol. 90, pp. 200–205, 2016.
14. A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet classification with deep convolutional neural networks," *Advances in Neural Information Processing Systems*, vol. 25, 2012.
15. M. Tan and Q. Le, "EfficientNet: Rethinking model scaling for convolutional neural networks," in *Proceedings of the International Conference on Machine Learning (ICML)*, 2019, pp. 6105–6114.
16. M. Tan and Q. Le, "EfficientNetV2: Smaller models and faster training," in *Proceedings of the International Conference on Machine Learning (ICML)*, 2021.
17. O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional networks for biomedical image segmentation," in *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, 2015, pp. 234–241.